

## Task-5:

### Sales Prediction Using Python:

Sales prediction means predicting how much of a product people will buy based on factors, a product and service-based business always need their Data Scientist to predict their future sales with every step they take to manipulate the cost of advertising their product.

by Mouli Nahal

```
In [2]: #Import required libraries

import numpy as np
import pandas as pd

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_error

import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [3]: df = pd.read_csv("Advertising.csv")
df
```

```
Out[3]:
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
...	...	...	...	...	...
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

200 rows × 5 columns

```
In [4]: df.shape
```

```
Out[4]: (200, 5)
```

In [5]: `df.columns`

Out[5]: Index(['Unnamed: 0', 'TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   200 non-null    int64
1   TV           200 non-null    float64
2   Radio        200 non-null    float64
3   Newspaper    200 non-null    float64
4   Sales        200 non-null    float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
```

In [7]: *#Checking gthe numbeer of missing values*

`df.isnull().sum()`

Out[7]: Unnamed: 0 0  
TV 0  
Radio 0  
Newspaper 0  
Sales 0  
dtype: int64

In [8]: `df.describe()`

Out[8]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
<b>count</b>	200.000000	200.000000	200.000000	200.000000	200.000000
<b>mean</b>	100.500000	147.042500	23.264000	30.554000	14.022500
<b>std</b>	57.879185	85.854236	14.846809	21.778621	5.217457
<b>min</b>	1.000000	0.700000	0.000000	0.300000	1.600000
<b>25%</b>	50.750000	74.375000	9.975000	12.750000	10.375000
<b>50%</b>	100.500000	149.750000	22.900000	25.750000	12.900000
<b>75%</b>	150.250000	218.825000	36.525000	45.100000	17.400000
<b>max</b>	200.000000	296.400000	49.600000	114.000000	27.000000

In [9]: `df.corr()`

Out[9]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
<b>Unnamed: 0</b>	1.000000	0.017715	-0.110680	-0.154944	-0.051616
<b>TV</b>	0.017715	1.000000	0.054809	0.056648	0.782224
<b>Radio</b>	-0.110680	0.054809	1.000000	0.354104	0.576223
<b>Newspaper</b>	-0.154944	0.056648	0.354104	1.000000	0.228299
<b>Sales</b>	-0.051616	0.782224	0.576223	0.228299	1.000000

```
In [10]: df.drop("Unnamed: 0",axis=1,inplace=True)
```

```
In [11]: x=df.drop("Sales",axis=1)
y=df["Sales"]
```

```
In [12]: x
```

```
Out[12]:
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4
...	...	...	...
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

200 rows × 3 columns

```
In [13]: y
```

```
Out[13]:
```

0	22.1
1	10.4
2	9.3
3	18.5
4	12.9
...	...
195	7.6
196	9.7
197	12.8
198	25.5
199	13.4

Name: Sales, Length: 200, dtype: float64

```
In [14]: x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_st
```

In [15]: `x_train`

Out[15]:

	TV	Radio	Newspaper
159	131.7	18.4	34.6
62	239.3	15.5	27.3
166	17.9	37.6	21.6
45	175.1	22.5	31.5
23	228.3	16.9	26.2
...	...	...	...
73	129.4	5.7	31.3
144	96.2	14.8	38.9
118	125.7	36.9	79.2
189	18.7	12.1	23.4
99	135.2	41.7	45.9

160 rows × 3 columns

In [16]: `model=LinearRegression()`  
`model`

Out[16]: `LinearRegression()`

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [17]: `model.fit(x_train,y_train)`

Out[17]: `LinearRegression()`

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [18]: `y_pred_test=model.predict(x_test)`  
`y_pred_test`

Out[18]: `array([ 6.72296139, 13.75631244, 12.04517877, 9.69967488, 15.5690964 ,`  
`10.05604087, 13.84462636, 11.57199727, 19.50132767, 10.63532913,`  
`18.86878088, 7.66094904, 14.95238223, 16.44076381, 21.16915592,`  
`20.34689481, 17.4717922 , 20.44501887, 6.6345129 , 11.59504468,`  
`10.73339912, 14.9550131 , 20.47462735, 14.4053483 , 19.82336217,`  
`21.23329569, 13.95616746, 23.25585899, 8.33169334, 20.72918361,`  
`15.86252542, 15.18601062, 21.74577286, 18.34562985, 12.80653699,`  
 `9.38349318, 5.37438571, 12.02365628, 19.26275296, 8.12044348])`

```
In [19]: y_test
```

```
Out[19]: 119      6.6
          77      14.2
          148     10.9
          149     10.1
          154     15.6
          151     11.6
          122     11.6
           6      11.8
          28      18.9
          71      12.4
          188     15.9
          34       9.5
          21      12.5
          40      16.6
          55      23.7
          104     20.7
           3      18.5
          39      21.5
          117      9.4
          134     10.8
          180     10.5
          26      15.0
          54      20.2
          165     11.9
          53      21.2
          93      22.2
          174     11.5
          17      24.4
          56       5.5
          84      21.7
          25      12.0
          156     15.3
          42      20.7
          141     19.2
          50      11.4
          51      10.7
          195      7.6
          116     12.2
          142     20.1
           24       9.7
          Name: Sales, dtype: float64
```

```
In [20]: MSE = mean_squared_error(y_test,y_pred_test)
          print("Mean squared is : ", MSE)
```

```
Mean squared is : 2.550964615953105
```

```
In [21]: MSE = mean_squared_error(y_pred_test,y_test)
          print("Mean squared is : ", MSE)
```

```
Mean squared is : 2.550964615953105
```

```
In [22]: RMSE = np.sqrt(MSE)
print("Root mean squared error : ", RMSE)
```

Root mean squared error : 1.5971739466799177

```
In [23]: mean_absolute_error(y_pred_test,y_test)
```

Out[23]: 1.2673937159929238

```
In [24]: r2score=r2_score(y_pred_test,y_test)
r2score
```


Out[24]: 0.8927421354788029

```
In [25]: r2score=r2_score(y_test,y_pred_test)
r2score
```

Out[25]: 0.8984204533332628

```
In [26]: plt.figure(figsize=(8,4))
sns.heatmap(df.corr(),annot=True)
```

Out[26]: <Axes: >



```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

